

```

import pandas as pd
import numpy as np
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, MinMaxScaler
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats

# Create a sample dataset
np.random.seed(42)
data = {
    'ID': range(1, 101),
    'Age': np.random.randint(20, 70, size=100),
    'Salary': np.append(np.random.randint(30000, 100000, size=95), [1000000, 1000000, -1000, 500, 600]),
    'Department': np.random.choice(['HR', 'Finance', 'IT', 'hr', 'finance', 'it'], size=100),
    'JoinDate': pd.date_range(start='1/1/2020', periods=100, freq='M'),
    'PerformanceScore': np.random.choice([1, 2, 3, 4, np.nan], size=100)
}

df = pd.DataFrame(data)

# Introduce some missing values
df.loc[5:10, 'Age'] = np.nan
df.loc[15:20, 'Salary'] = np.nan
df.loc[25:30, 'Department'] = np.nan

# Initial Inspection
print("Initial Dataset Info:")
print(df.info())
print("\nInitial Dataset Description:")
print(df.describe())
print("\nFirst Few Rows of the Initial Dataset:")
print(df.head())

# Handling Missing Values
print("\nHandling Missing Values:")
# Impute numerical columns with mean
num_cols = df.select_dtypes(include=np.number).columns

```

0s completed at 1:08 PM

```
+ Code + Text All changes saved

# Handling Missing Values
print("\nHandling Missing Values:")
# Impute numerical columns with mean
num_cols = df.select_dtypes(include=np.number).columns
imputer_num = SimpleImputer(strategy='mean')
df[num_cols] = imputer_num.fit_transform(df[num_cols])

# Impute categorical columns with most frequent value
cat_cols = df.select_dtypes(include='object').columns
imputer_cat = SimpleImputer(strategy='most_frequent')
df[cat_cols] = imputer_cat.fit_transform(df[cat_cols])

print("\nMissing Values After Imputation:")
print(df.isnull().sum())

# Outlier Detection and Treatment
print("\nDetecting and Handling Outliers:")
z_scores = np.abs(stats.zscore(df[num_cols]))
outliers = np.where(z_scores > 3, True, False)
df_no_outliers = df[~outliers.any(axis=1)]

print("\nDataset Shape Before Removing Outliers:", df.shape)
print("Dataset Shape After Removing Outliers:", df_no_outliers.shape)

# Normalization/Standardization
print("\nNormalizing/Standardizing Numerical Columns:")
scaler = MinMaxScaler()
df_scaled = df_no_outliers.copy()
df_scaled[num_cols] = scaler.fit_transform(df_scaled[num_cols])

# Handling Inconsistencies in Categorical Data
print("\nHandling Inconsistencies in Categorical Data:")
df_scaled['Department'] = df_scaled['Department'].str.strip().str.lower()
df_scaled['Department'] = df_scaled['Department'].replace({'hr': 'HR', 'finance': 'Finance', 'it': 'IT'})

print("\nFinal Cleaned Dataset Info:")
print(df_scaled.info())
print("\nFinal Cleaned Dataset Description:")
print(df_scaled.describe())

0s completed at 1:08 PM
```

```
+ Code + Text All changes saved
df_scaled['Department'] = df_scaled['Department'].str.strip().str.lower()
df_scaled['Department'] = df_scaled['Department'].replace({'hr': 'HR', 'finance': 'Finance', 'it': 'IT'})

print("\nFinal Cleaned Dataset Info:")
print(df_scaled.info())
print("\nFinal Cleaned Dataset Description:")
print(df_scaled.describe())
print("\nFirst Few Rows of the Final Cleaned Dataset:")
print(df_scaled.head())

# Save the cleaned dataset
df_scaled.to_csv('cleaned_data.csv', index=False)
print("\nCleaned dataset saved to 'cleaned_data.csv'")
```

Initial Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 6 columns):
Column Non-Null Count Dtype
0 ID 100 non-null int64
1 Age 94 non-null float64
2 Salary 94 non-null float64
3 Department 94 non-null object
4 JoinDate 100 non-null datetime64[ns]
5 PerformanceScore 80 non-null float64
dtypes: datetime64[ns](1), float64(3), int64(1), object(1)
memory usage: 4.8+ KB
None

Initial Dataset Description:

	ID	Age	Salary	JoinDate	PerformanceScore
count	100.000000	94.000000	94.000000	100	80.000000
mean	50.500000	44.351064	85033.787234	2024-03-15 18:28:48	75.000000
min	1.000000	20.000000	-1000.000000	2020-01-31 00:00:00	50.000000
25%	25.750000	33.000000	48070.500000	2022-02-21 00:00:00	60.000000
50%	50.500000	44.000000	69006.500000	2024-03-15 12:00:00	70.000000
75%	75.250000	58.000000	85785.000000	2026-04-07 12:00:00	80.000000
max	100.000000	69.000000	100000.000000	2028-04-30 00:00:00	90.000000
std	29.011492	14.668163	137762.627188	NaN	10.000000

completed at 1:08 PM

```
+ Code + Text All changes saved
0s print("\nCleaned dataset saved to 'cleaned_data.csv'")

5 PerformanceScore 80 non-null float64
dtypes: datetime64[ns](1), float64(3), int64(1), object(1)
memory usage: 4.8+ KB
None

Initial Dataset Description:
count ID Age Salary JoinDate \
mean 50.500000 44.351064 85033.787234 2024-03-15 18:28:48
min 1.000000 20.000000 -1000.000000 2020-01-31 00:00:00
25% 25.750000 33.000000 48070.500000 2022-02-21 00:00:00
50% 50.500000 44.000000 69006.500000 2024-03-15 12:00:00
75% 75.250000 58.000000 85785.000000 2026-04-07 12:00:00
max 100.000000 69.000000 100000.000000 2028-04-30 00:00:00
std 29.011492 14.668163 137762.627188 NaN

PerformanceScore
count 80.000000
mean 2.525000
min 1.000000
25% 2.000000
50% 2.000000
75% 4.000000
max 4.000000
std 1.147093

First Few Rows of the Initial Dataset:
ID Age Salary Department JoinDate PerformanceScore
0 1 58.0 32568.0 Finance 2020-01-31 NaN
1 2 48.0 92592.0 it 2020-02-29 1.0
2 3 34.0 97563.0 it 2020-03-31 4.0
3 4 62.0 32695.0 IT 2020-04-30 1.0
4 5 27.0 78190.0 HR 2020-05-31 NaN

Handling Missing Values:

Missing Values After Imputation:
ID 0
Age 0
Salary 0
Department 0
JoinDate 0
```

0s completed at 1:08 PM

```
+ Code + Text All changes saved
print("\nCleaned dataset saved to 'cleaned_data.csv'")

Handling Missing Values:
Missing Values After Imputation:
ID      0
Age     0
Salary  0
Department 0
JoinDate 0
PerformanceScore 0
dtype: int64

Detecting and Handling Outliers:
Dataset Shape Before Removing Outliers: (100, 6)
Dataset Shape After Removing Outliers: (98, 6)

Normalizing/Standardizing Numerical Columns:

Handling Inconsistencies in Categorical Data:

Final Cleaned Dataset Info:
<class 'pandas.core.frame.DataFrame'>
Index: 98 entries, 0 to 99
Data columns (total 6 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   ID                   98 non-null    float64
1   Age                  98 non-null    float64
2   Salary               98 non-null    float64
3   Department           98 non-null    object
4   JoinDate             98 non-null    datetime64[ns]
5   PerformanceScore     98 non-null    float64
dtypes: datetime64[ns](1), float64(4), object(1)
memory usage: 5.4+ KB
None

Final Cleaned Dataset Description:
      ID      Age      Salary      JoinDate \
count  98.000000  98.000000  98.000000      98
mean    0.490517  0.495024  0.672514  2024-02-16 04:24:29.387755008
min     0.000000  0.000000  0.000000  2020-01-31 00:00:00
max     0.999999  0.999999  0.999999  2024-02-16 04:24:29.387755008

0s completed at 1:08 PM
```

```
+ Code + Text All changes saved
1 Age 98 non-null float64
2 Salary 98 non-null float64
3 Department 98 non-null object
4 JoinDate 98 non-null datetime64[ns]
5 PerformanceScore 98 non-null float64
dtypes: datetime64[ns](1), float64(4), object(1)
memory usage: 5.4+ KB
None

Final Cleaned Dataset Description:
      ID      Age      Salary      JoinDate \
count 98.000000 98.000000 98.000000 98
mean  0.490517 0.495024 0.672514 2024-02-16 04:24:29.387755008
min    0.000000 0.000000 0.000000 2020-01-31 00:00:00
25%    0.244949 0.270408 0.499828 2022-02-07 00:00:00
50%    0.489899 0.493378 0.719133 2024-02-14 12:00:00
75%    0.734848 0.750000 0.858938 2026-02-21 00:00:00
max    1.000000 1.000000 1.000000 2028-04-30 00:00:00
std    0.288277 0.289471 0.240970 NaN

PerformanceScore
count 98.000000
mean  0.505102
min    0.000000
25%    0.333333
50%    0.508333
75%    0.666667
max    1.000000
std    0.340960

First Few Rows of the Final Cleaned Dataset:
      ID      Age      Salary Department  JoinDate  PerformanceScore
0  0.000000 0.775510 0.335134    Finance 2020-01-31 0.508333
1  0.010101 0.571429 0.934397         IT 2020-02-29 0.000000
2  0.020202 0.285714 0.984026         IT 2020-03-31 1.000000
3  0.030303 0.857143 0.336402         IT 2020-04-30 0.000000
4  0.040404 0.142857 0.798611         HR 2020-05-31 0.508333

Cleaned dataset saved to 'cleaned_data.csv'
```

0s completed at 1:08 PM


```
+ Code + Text All changes saved
pip install pandas numpy statsmodels matplotlib

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.0.3)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.25.2)
Requirement already satisfied: statsmodels in /usr/local/lib/python3.10/dist-packages (0.14.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.1)
Requirement already satisfied: scipy!=1.9.2,>=1.8 in /usr/local/lib/python3.10/dist-packages (from statsmodels) (1.11.4)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.10/dist-packages (from statsmodels) (0.5.6)
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.10/dist-packages (from statsmodels) (24.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.53.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.2)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy>=0.5.6->statsmodels) (1.16.0)

[2] import pandas as pd
import numpy as np

# Generate synthetic stock prices
np.random.seed(42)
dates = pd.date_range(start='2020-01-01', end='2023-01-01', freq='B')
price = np.random.normal(loc=0.001, scale=0.02, size=len(dates))
price[0] = 0
price = 100 + np.cumsum(price)

stock_data = pd.DataFrame({'Date': dates, 'Price': price})
stock_data.set_index('Date', inplace=True)

[3] import matplotlib.pyplot as plt

plt.figure(figsize=(12, 6))
plt.plot(stock_data.index, stock_data['Price'], label='Stock Price')
```

0s completed at 5:41 PM




```
+ Code + Text All changes saved
from statsmodels.tsa.arima.model import ARIMA

# Define the ARIMA model
model = ARIMA(stock_data['Price'], order=(1, 1, 1))
model_fit = model.fit()

# Print summary of the model
print(model_fit.summary())
```

/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency B
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency B
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency B
self._init_dates(dates, freq)

SARIMAX Results

Dep. Variable:	Price	No. Observations:	783
Model:	ARIMA(1, 1, 1)	Log Likelihood	1959.018
Date:	Thu, 20 Jun 2024	AIC	-3912.037
Time:	12:10:59	BIC	-3898.051
Sample:	01-01-2020 - 12-30-2022	HQIC	-3906.658
Covariance Type:	opg		

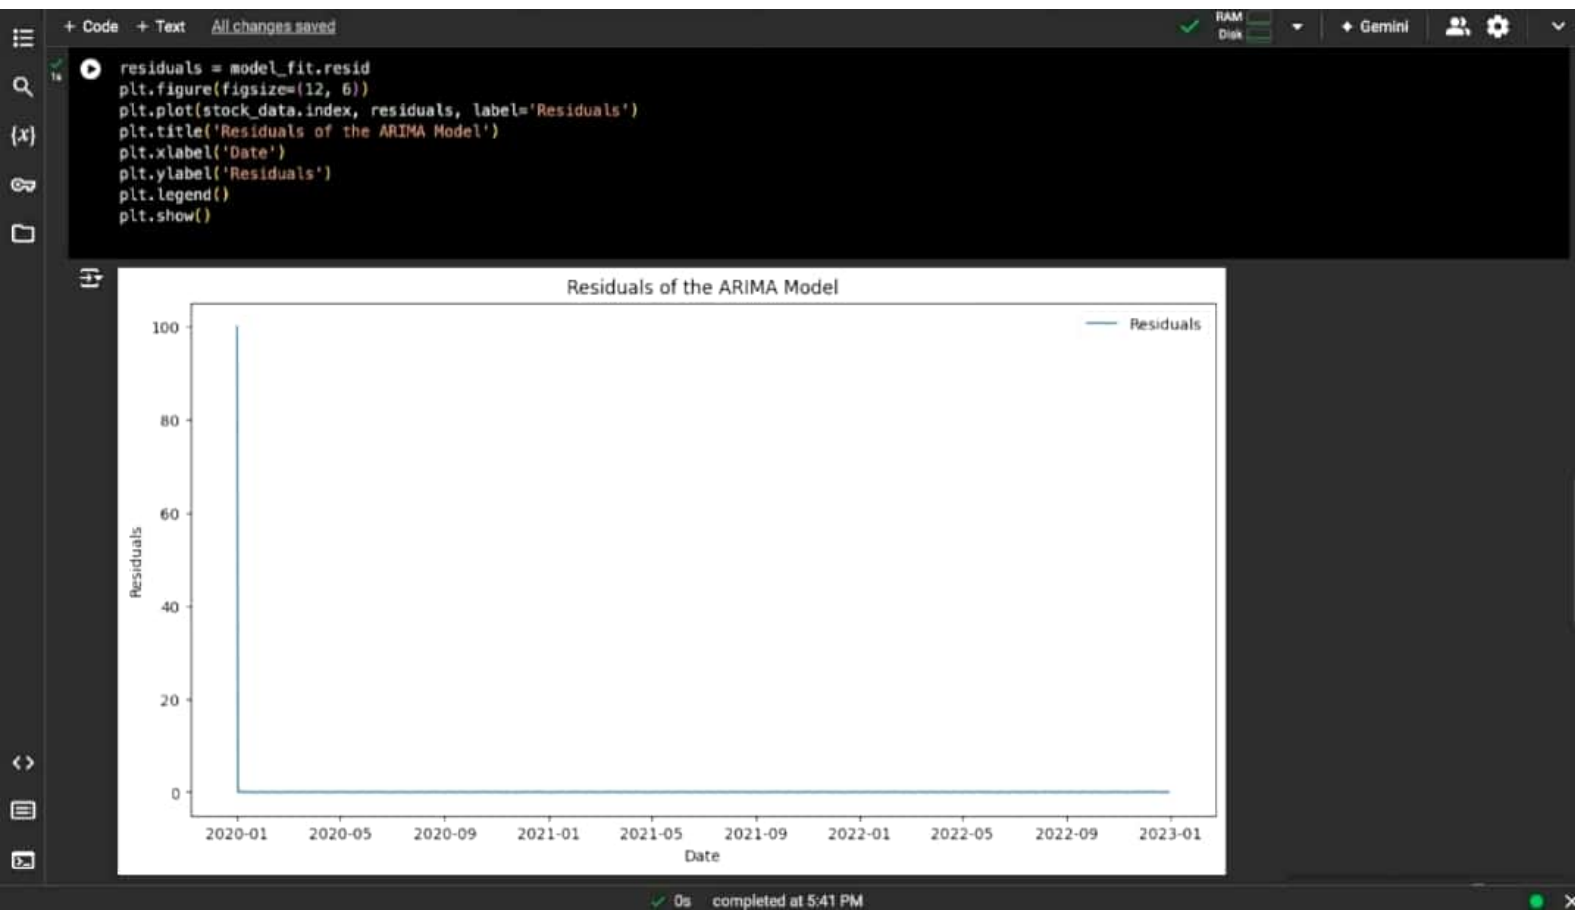
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.0824	1.568	-0.053	0.958	-3.156	2.992
ma.L1	0.0587	1.571	0.037	0.970	-3.020	3.138
sigma2	0.0004	1.92e-05	20.311	0.000	0.000	0.000

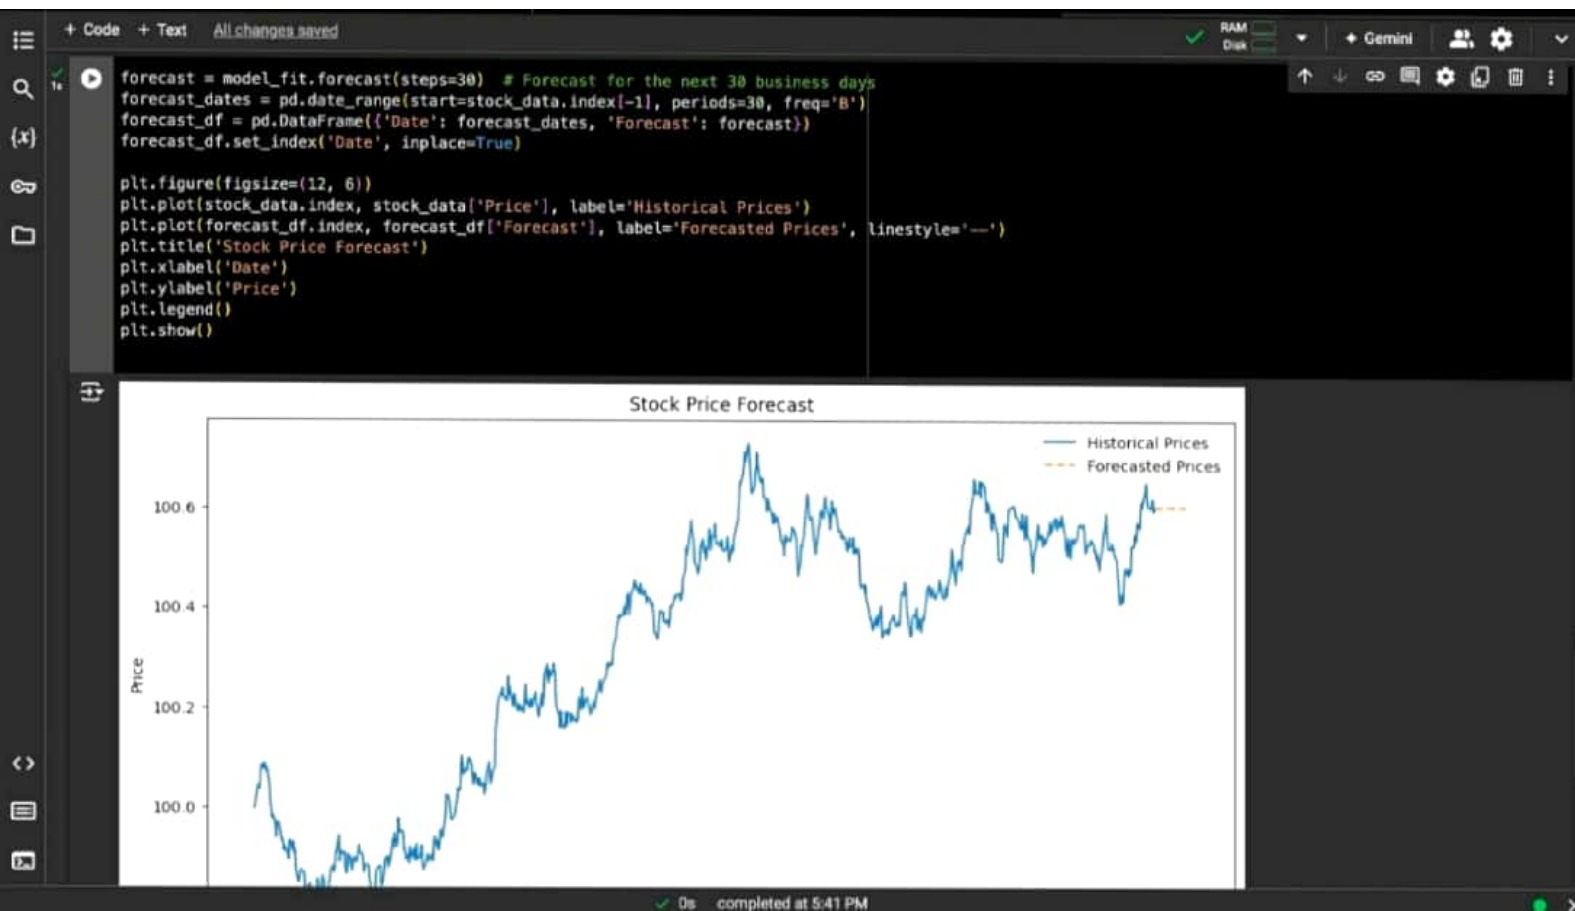
Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):	3.39
Prob(Q):	1.00	Prob(JB):	0.18
Heteroskedasticity (H):	1.08	Skew:	0.15
Prob(H) (two-sided):	0.55	Kurtosis:	3.13

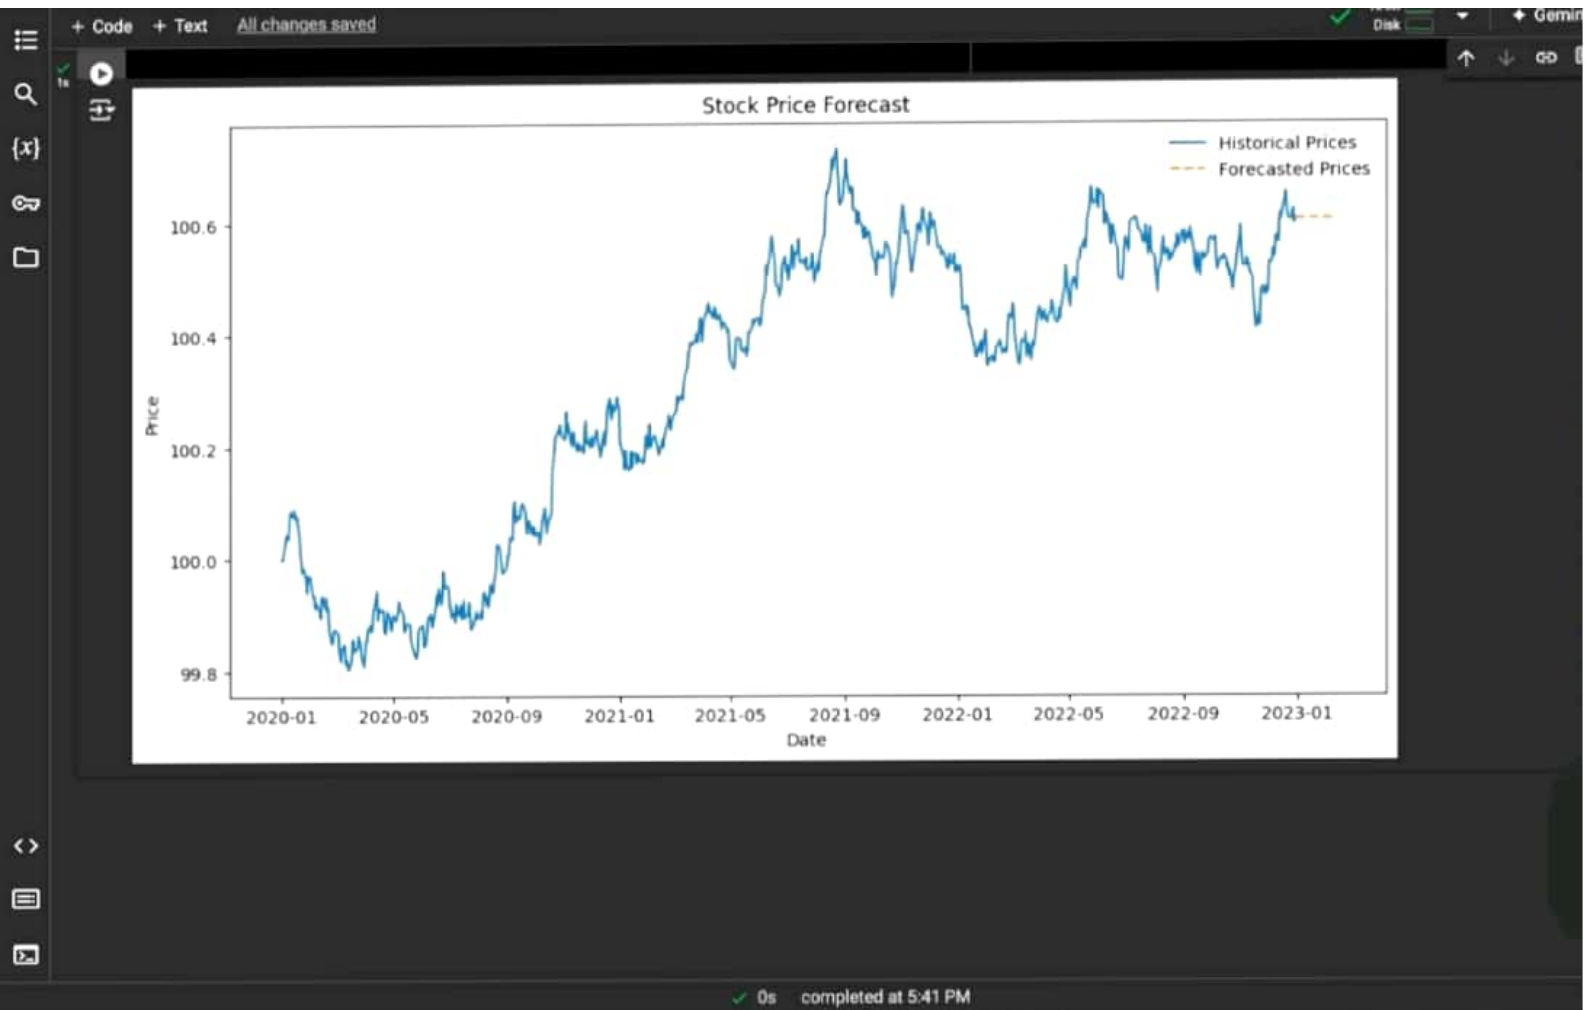
Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

0s completed at 5:41 PM







```
+ Code + Text All changes saved
import pandas as pd
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori, association_rules

# Sample transactional data
data = {
    'TransactionID': [1, 2, 3, 4, 5],
    'Items': [
        ['milk', 'bread', 'butter'],
        ['bread', 'butter', 'eggs'],
        ['milk', 'bread', 'eggs'],
        ['milk', 'butter', 'eggs'],
        ['bread', 'butter']
    ]
}

# Convert to DataFrame
df = pd.DataFrame(data)

# Extract list of transactions
transactions = df['Items'].tolist()

# Initialize transaction encoder
te = TransactionEncoder()

# Transform the list of transactions into a one-hot encoded DataFrame
te_ary = te.fit(transactions).transform(transactions)
df_transformed = pd.DataFrame(te_ary, columns=te.columns_)

# Apply Apriori algorithm
frequent_itemsets = apriori(df_transformed, min_support=0.6, use_colnames=True)

# Generate association rules
rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1)

print(frequent_itemsets)
print(rules)
```

0s completed at 1:56 PM

```
+ Code + Text All changes saved
# Transform the list of transactions into a one-hot encoded DataFrame
te_ary = te.fit(transactions).transform(transactions)
df_transformed = pd.DataFrame(te_ary, columns=te.columns_)

# Apply Apriori algorithm
frequent_itemsets = apriori(df_transformed, min_support=0.6, use_colnames=True)

# Generate association rules
rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1)

print(frequent_itemsets)
print(rules)
```

	support	itemsets
0	0.8	(bread)
1	0.8	(butter)
2	0.6	(eggs)
3	0.6	(milk)
4	0.6	(butter, bread)

Empty DataFrame
Columns: [antecedents, consequents, antecedent support, consequent support, support, confidence, lift, leverage, conviction, zhangs_metric]
Index: []


```
+ Code + Text All changes saved
import pandas as pd
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split

# Sample dataset creation
data = {
    'customerID': ['C001', 'C002', 'C003', 'C004', 'C005', 'C006', 'C007', 'C008', 'C009', 'C010'],
    'gender': ['Female', 'Male', 'Female', 'Male', 'Male', 'Female', 'Female', 'Male', 'Male', 'Female'],
    'SeniorCitizen': [0, 1, 0, 0, 1, 1, 0, 0, 0, 1],
    'Partner': ['Yes', 'No', 'Yes', 'No', 'No', 'Yes', 'No', 'No', 'Yes', 'Yes'],
    'Dependents': ['No', 'No', 'Yes', 'No', 'No', 'No', 'Yes', 'No', 'Yes', 'No'],
    'tenure': [1, 34, 2, 45, 5, 2, 8, 22, 10, 30],
    'PhoneService': ['No', 'Yes', 'Yes', 'No', 'Yes', 'Yes', 'No', 'Yes', 'Yes', 'Yes'],
    'InternetService': ['DSL', 'Fiber optic', 'DSL', 'Fiber optic', 'DSL', 'DSL', 'Fiber optic', 'DSL', 'DSL', 'Fiber optic'],
    'Contract': ['Month-to-month', 'One year', 'Month-to-month', 'One year', 'Month-to-month', 'Month-to-month', 'One year', 'Month-to-month', 'Month-to-month', 'Month-to-month'],
    'MonthlyCharges': [29.85, 56.95, 53.85, 42.30, 70.70, 53.85, 99.65, 89.10, 29.75, 49.95],
    'TotalCharges': [29.85, 1889.50, 108.15, 1840.75, 151.65, 108.15, 820.50, 1949.40, 301.90, 1490.75],
    'Churn': ['Yes', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No']
}

# Create DataFrame
df = pd.DataFrame(data)

# Encode categorical variables
df_encoded = df.copy()
label_encoders = {}
for column in ['gender', 'Partner', 'Dependents', 'PhoneService', 'InternetService', 'Contract', 'Churn']:
    le = LabelEncoder()
    df_encoded[column] = le.fit_transform(df_encoded[column])
    label_encoders[column] = le

df_encoded.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	InternetService	Contract	MonthlyCharges	TotalCharges	Churn
0	C001	0	0	1	0	1	0	0	0	29.85	29.85	1

0s completed at 12:19 PM

+ Code
+ Text
All changes saved

RAM
Disk

Gemini

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	InternetService	Contract	MonthlyCharges	TotalCharges	Churn
0	C001	0	0	1	0	1	0	0	0	29.85	29.85	1
1	C002	1	1	0	0	34	1	1	1	56.95	1889.50	0
2	C003	0	0	1	1	2	1	0	0	53.85	108.15	1
3	C004	1	0	0	0	45	0	1	1	42.30	1840.75	0
4	C005	1	1	0	0	5	1	0	0	70.70	151.85	1

Next steps:
Generate code with df_encoded
View recommended plots

```

[1] # Define features and target variable
X = df_encoded.drop(['customerID', 'Churn'], axis=1)
y = df_encoded['Churn']

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

[3] scaler = StandardScaler()
X_train[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.fit_transform(X_train[['tenure', 'MonthlyCharges', 'TotalCharges']])
X_test[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.transform(X_test[['tenure', 'MonthlyCharges', 'TotalCharges']])

[4] from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

# Build and train the logistic regression model
logreg = LogisticRegression(random_state=42)
logreg.fit(X_train, y_train)

# Predict on test data
y_pred_logreg = logreg.predict(X_test)

```

0s
completed at 12:19 PM

+ Code + Text All changes saved

✓ RAM Disk Gemini

```
# Evaluate the logistic regression model
logreg_accuracy = accuracy_score(y_test, y_pred_logreg)
logreg_conf_matrix = confusion_matrix(y_test, y_pred_logreg)
logreg_class_report = classification_report(y_test, y_pred_logreg)

print("Logistic Regression Accuracy:", logreg_accuracy)
print("Confusion Matrix:\n", logreg_conf_matrix)
print("Classification Report:\n", logreg_class_report)
```

Logistic Regression Accuracy: 0.5
Confusion Matrix:
[[1 1]
 [0 0]]
Classification Report:

	precision	recall	f1-score	support
0	1.00	0.50	0.67	2
1	0.00	0.00	0.00	0
accuracy			0.50	2
macro avg	0.50	0.25	0.33	2
weighted avg	1.00	0.50	0.67	2

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to nan
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to nan
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to nan
_warn_prf(average, modifier, msg_start, len(result))

```
[5] from sklearn.tree import DecisionTreeClassifier

# Build and train the decision tree model
tree = DecisionTreeClassifier(random_state=42)
tree.fit(X_train, y_train)

# Predict on test data
y_pred_tree = tree.predict(X_test)
```

0s completed at 12:19 PM

Scanned with CamScanner

```
+ Code + Text All changes saved
✓ RAM
Disk
+ Gemini
1m
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import string
import gensim.downloader as api
from nltk.sentiment.vader import SentimentIntensityAnalyzer
import pandas as pd

# Sample dataset
data = {
    "review": [
        "I love this product! It's absolutely amazing.",
        "This is the worst thing I have ever bought. Totally useless.",
        "Not bad, could be better.",
        "Pretty decent product for the price.",
        "Absolutely horrible! Do not buy this.",
        "Fantastic quality and great value for money.",
        "Mediocre, nothing special.",
        "Exceeded my expectations. Highly recommend!",
        "Terrible experience, will never purchase again.",
        "Good, but not great. Satisfied overall."
    ]
}

reviews_df = pd.DataFrame(data)

# Download NLTK data
nltk.download('punkt')
nltk.download('stopwords')
nltk.download('vader_lexicon')

def clean_tokenize(text):
    text = text.lower()
    text = text.translate(str.maketrans('', '', string.punctuation))
    tokens = word_tokenize(text)
    tokens = [word for word in tokens if word not in stopwords.words('english')]
    return tokens

✓ 1m 10s completed at 3:24 PM
```

```
+ Code + Text All changes saved
✓ RAM
Disk
+ Gemini
1m
# Apply cleaning and tokenization
reviews_df['tokens'] = reviews_df['review'].apply(clean_tokenize)

# Load pre-trained word vectors
glove_vectors = api.load("glove-wiki-gigaword-50")

def get_word_vectors(tokens):
    vectors = [glove_vectors[word] for word in tokens if word in glove_vectors]
    return vectors

# Apply word embedding
reviews_df['word_vectors'] = reviews_df['tokens'].apply(get_word_vectors)

# Initialize VADER sentiment analyzer
sia = SentimentIntensityAnalyzer()

def get_sentiment_score(text):
    sentiment = sia.polarity_scores(text)
    return sentiment

# Apply sentiment analysis
reviews_df['sentiment'] = reviews_df['review'].apply(get_sentiment_score)

# Extract compound sentiment score
reviews_df['compound'] = reviews_df['sentiment'].apply(lambda x: x['compound'])

# Determine overall sentiment
reviews_df['sentiment_label'] = reviews_df['compound'].apply(lambda x: 'positive' if x > 0 else 'negative' if x < 0 else 'neutral')

# Aggregate results
sentiment_summary = reviews_df['sentiment_label'].value_counts()
reviews_df, sentiment_summary

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package vader_lexicon to /root/nltk_data...
1 100.0% 66.0/66.0MB downloaded
✓ 1m 10s completed at 3:24 PM
```



```
+ Code + Text All changes saved
[ nltk_data ] Downloading package stopwords to /root/nltk_data...
[ nltk_data ] Unzipping corpora/stopwords.zip.
[ nltk_data ] Downloading package vader_lexicon to /root/nltk_data...
[ ] 100.0% 66.0/66.0MB downloaded

( review \
0 I love this product! It's absolutely amazing.
1 This is the worst thing I have ever bought. To...
2 Not bad, could be better.
3 Pretty decent product for the price.
4 Absolutely horrible! Do not buy this.
5 Fantastic quality and great value for money.
6 Mediocre, nothing special.
7 Exceeded my expectations. Highly recommend!
8 Terrible experience, will never purchase again.
9 Good, but not great. Satisfied overall.

tokens \
0 [love, product, absolutely, amazing]
1 [worst, thing, ever, bought, totally, useless]
2 [bad, could, better]
3 [pretty, decent, product, price]
4 [absolutely, horrible, buy]
5 [fantastic, quality, great, value, money]
6 [mediocre, nothing, special]
7 [exceeded, expectations, highly, recommend]
8 [terrible, experience, never, purchase]
9 [good, great, satisfied, overall]

word_vectors \
0 [[-0.13886, 1.1401, -0.85212, -0.29212, 0.7553...
1 [[-0.14968, -0.42252, 0.16736, 0.13474, -0.310...
2 [[-0.17981, -0.40407, -0.1653, -0.60687, -0.39...
3 [[-0.24922, -0.39835, -0.45851, -0.34846, 0.74...
4 [[0.36582, -0.43975, -0.35816, 0.096443, 0.995...
5 [[0.3333, 0.30612, -0.63572, 0.051507, 0.78602...
6 [[-1.0406, -0.61579, -0.28125, -0.51557, 0.079...
7 [[-0.32462, -0.079688, 1.2704, -0.55724, 0.029...
8 [[0.33209, -0.028359, -0.58145, -0.4487, 0.254...
9 [[-0.35586, 0.5213, -0.6107, -0.30131, 0.94862...

sentiment compound sentiment_label
0 {'neg': 0.0, 'neu': 0.318, 'pos': 0.682, 'comp... 0.8620 positive
1 {'neg': 0.473, 'neu': 0.527, 'pos': 0.0, 'comp... -0.8016 negative
2 {'neg': 0.0, 'neu': 0.343, 'pos': 0.657, 'comp... 0.6956 positive

1m 10s completed at 3:24 PM
```

```
+ Code + Text All changes saved
8 Terrible experience, will never purchase again.
9 Good, but not great. Satisfied overall.

tokens \
0 [love, product, absolutely, amazing]
1 [worst, thing, ever, bought, totally, useless]
2 [bad, could, better]
3 [pretty, decent, product, price]
4 [absolutely, horrible, buy]
5 [fantastic, quality, great, value, money]
6 [mediocre, nothing, special]
7 [exceeded, expectations, highly, recommend]
8 [terrible, experience, never, purchase]
9 [good, great, satisfied, overall]

word_vectors \
0 [[-0.13886, 1.1401, -0.85212, -0.29212, 0.7553...
1 [[-0.14968, -0.42252, 0.16736, 0.13474, -0.310...
2 [[-0.17981, -0.40407, -0.1653, -0.60687, -0.39...
3 [[-0.24922, -0.39835, -0.45851, -0.34846, 0.74...
4 [[0.36582, -0.43975, -0.35816, 0.096443, 0.995...
5 [[0.3333, 0.30612, -0.63572, 0.051507, 0.78602...
6 [[-1.0406, -0.61579, -0.28125, -0.51557, 0.079...
7 [[-0.32462, -0.079688, 1.2704, -0.55724, 0.029...
8 [[0.33209, -0.028359, -0.58145, -0.4487, 0.254...
9 [[-0.35586, 0.5213, -0.6107, -0.30131, 0.94862...

sentiment compound sentiment_label
0 {'neg': 0.0, 'neu': 0.318, 'pos': 0.682, 'comp... 0.8620 positive
1 {'neg': 0.473, 'neu': 0.527, 'pos': 0.0, 'comp... -0.8016 negative
2 {'neg': 0.0, 'neu': 0.343, 'pos': 0.657, 'comp... 0.6956 positive
3 {'neg': 0.0, 'neu': 0.61, 'pos': 0.39, 'compou... 0.4939 positive
4 {'neg': 0.45, 'neu': 0.55, 'pos': 0.0, 'compou... -0.6230 negative
5 {'neg': 0.0, 'neu': 0.284, 'pos': 0.716, 'comp... 0.8779 positive
6 {'neg': 0.53, 'neu': 0.47, 'pos': 0.0, 'compou... -0.3089 negative
7 {'neg': 0.0, 'neu': 0.565, 'pos': 0.435, 'comp... 0.4740 positive
8 {'neg': 0.383, 'neu': 0.617, 'pos': 0.0, 'comp... -0.4767 negative
9 {'neg': 0.6, 'neu': 0.242, 'pos': 0.157, 'comp... -0.7571 negative
sentiment_label
positive 5
negative 5
Name: count, dtype: int64

1m 10s completed at 3:24 PM
```